

## TensorFlow skill testing project: Classification in TensorFlow

**Objective:** Create and run your own machine learning algorithms in TensorFlow on a real-world dataset.

**Architecture:** Your code will consist of three modules. This document contains the specification for the second module.

---

### Module 2: `graph_constructor.py`

You have complete freedom to design this module any way that makes sense to you. Your goal will be to build three different computational graphs. The first will be a simple logistic regression model, the second will be a two-layer neural net, and the third will be a  $k$ -nearest-neighbour classifier. These graphs will be used to train and test your algorithm using real world data. This module, as well as `data_preprocessing.py`, are essentially helper functions for `learn.py`, which you'll build out next.

---

### Common mistakes

There are a few things to keep an eye on when you start working with TensorFlow. Here are some things you should watch out for as you build this module:

1. **Training ops:** Training ops are functions like

```
tf.train.GradientDescentOptimizer(learning_rate).minimize(loss_function)
```

In order to make your algorithm learn, you'll have to *explicitly run these* in a TensorFlow session. In other words, if your training op is called `train_op()`, you'll need a line like

```
train_op_val = session.run(train_op, feed_dict=feed_dict)
```

in your code.

2. **Cross-entropy:** the function

```
tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=labels)
```

takes **logits** as one of its arguments. **logits** is supposed to be the raw output of your model *prior to normalization* by the **softmax()** function. Be sure not to take the **softmax()** of your logits *prior* to feeding them into the cross-entropy function, because **softmax\_cross\_entropy\_with\_logits()** implements the **softmax()** itself internally.